

btc-stats-models

```
btc_df <- read_csv("data/BTC-Daily.csv")
```

```
## Rows: 2651 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr  (1): symbol
## dbl  (7): unix, open, high, low, close, Volume BTC, Volume USD
## dtm  (1): date
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
x = zoo(btc_df$close, as.Date(btc_df$date))
```

```
log_returns = diff(log(x))
```


```
head(log_returns)
```

```
##      2014-11-29      2014-11-30      2014-12-01      2014-12-02      2014-12-03      2014-12-04
## 0.001168659 -0.009012673  0.013435877  0.002270208 -0.006826146 -0.030924324
```

```
tail(log_returns)
```


```
##      2022-02-24      2022-02-25      2022-02-26      2022-02-27      2022-02-28
## 0.0291543332 0.0220283692 -0.0021684582 -0.0373187285  0.1353574253
##      2022-03-01
## 0.0001505249
```

```
plot(log_returns)
```




```
acf(coredata(log_returns), main="Sample Autocorrelation of Daily Log-Returns")
```


Sample Autocorrelation of Daily Log-Returns



```
plot(coredata(x))
```




```
abs_lrs <- abs(log_returns)
plot(abs_lrs)
```




```
acf(coredata(abs_lrs), main="Sample Autocorrelation of Daily Absolute Log>Returns")
```

Sample Autocorrelation of Daily Absolute Log-Returns




```
qqnorm(log_returns, main="Q-Q Plot for Daily Log-Returns")  
qqline(log_returns, col="red")
```

Q-Q Plot for Daily Log-Returns



```
qqnorm(abs_lrs, main="Q-Q Plot for Daily Log-Returns")  
qqline(abs_lrs, col="red")
```

Q-Q Plot for Daily Log-Returns



Split into train, test

```
n <- length(log_returns)

# Split index
test_size <- 30
train <- log_returns[1:(n - test_size)]
test <- log_returns[(n - test_size + 1):n]
```

Arima model

```
fit_auto <- auto.arima(train)
summary(fit_auto)

## Series: train
## ARIMA(0,0,1) with non-zero mean
##
## Coefficients:
##          ma1      mean
##       -0.0513  0.0018
## s.e.    0.0192  0.0007
##
## sigma^2 = 0.001612:  log likelihood = 4707.26
```



```
## AIC=-9408.52   AICc=-9408.51   BIC=-9390.91
##
## Training set error measures:
##           ME           RMSE           MAE MPE MAPE           MASE           ACF1
## Training set 1.38876e-07 0.0401311 0.02591193 NaN  Inf 0.658159 -0.0009093721
```

```
forecast_auto <- forecast(fit_auto, h = 30)
```

ARIMA(0,0,1) with non-zero mean

Coefficients: mal mean -0.0513 0.0018 s.e. 0.0192 0.0007

```
library(forecast)

fit <- auto.arima(train)
fc <- forecast(fit, h = test_size)

accuracy(fc, test)
```

```
##           ME           RMSE           MAE           MPE           MAPE           MASE
## Training set 1.388760e-07 0.04013110 0.02591193           NaN           Inf 0.6581590
## Test set      2.540786e-03 0.04155324 0.02821374 62.67331 133.9955 0.7166247
##           ACF1
## Training set -0.0009093721
## Test set      NA
```

metrics in ordinary format (not log)

```
# Get the last known price before forecast starts
last_price <- coredata(x)[length(x) - test_size]

# Convert forecasted log-returns to prices
predicted_prices <- numeric(test_size)
predicted_prices[1] <- last_price * exp(fc$mean[1])

for (i in 2:test_size) {
  predicted_prices[i] <- predicted_prices[i - 1] * exp(fc$mean[i])
}

# Get true prices to compare against
true_prices <- coredata(x)[(length(x) - test_size + 1):(length(x))]

# Evaluate metrics manually
mae <- mean(abs(true_prices - predicted_prices))
rmse <- sqrt(mean((true_prices - predicted_prices)^2))

cat("MAE: ", round(mae, 4), "\n")
```


```
## MAE: 2671.279
```

```
cat("RMSE:", round(rmse, 4), "\n")
```

```
## RMSE: 3167.454
```

```
arima_pred_prices = predicted_prices
```

```
plot(index(x)[(length(x) - test_size + 1):length(x)], true_prices,
     type = "l", col = "black", ylab = "Price", xlab = "Date", main = "Forecast vs Actual")
lines(index(x)[(length(x) - test_size + 1):length(x)], predicted_prices, col = "red")
legend("topright", legend = c("Actual", "Forecast"), col = c("black", "red"), lty = 1)
```



Exponential smoothing


```
library(forecast)
fit_ets <- ets(train) # train is already in log scale.
fc_ets <- forecast(fit_ets, h = test_size)

accuracy(fc_ets)
```

```
##               ME          RMSE          MAE   MPE  MAPE          MASE         ACF1
## Training set 3.584627e-06 0.04018761 0.02601443 -Inf   Inf  0.6607625 -0.05288728
```

```
actual_log_returns <- test
forecasted_log_returns <- fc_ets$mean
```

```
# Plot
plot(actual_log_returns, type = "l", col = "black",
      ylab = "Log-Return", xlab = "Time", main = "Forecasted vs Actual Log-Returns")
lines(forecasted_log_returns, col = "blue")
legend("topright", legend = c("Actual", "Forecast"), col = c("black", "blue"), lty = 1)
```



```
last_price <- coredata(x)[length(x) - test_size] # last actual price before forecast
```

```
# Reconstruct price forecast
```

```
predicted_prices_ets <- numeric(test_size)
```

```
predicted_prices_ets[1] <- last_price * exp(fc_ets$mean[1])
```

```
for (i in 2:test_size) {
```

```
  predicted_prices_ets[i] <- predicted_prices_ets[i - 1] * exp(fc_ets$mean[i])
```

```
}
```

```
true_prices <- coredata(x)[(length(x) - test_size + 1):(length(x))]
```

```
# Metrics
```

```
mae_ets <- mean(abs(true_prices - predicted_prices_ets))
```

```
rmse_ets <- sqrt(mean((true_prices - predicted_prices_ets)^2))
```

```
cat("ETS Forecast\n")
```

```
## ETS Forecast
```

```
cat("MAE: ", round(mae_ets, 4), "\n")
```

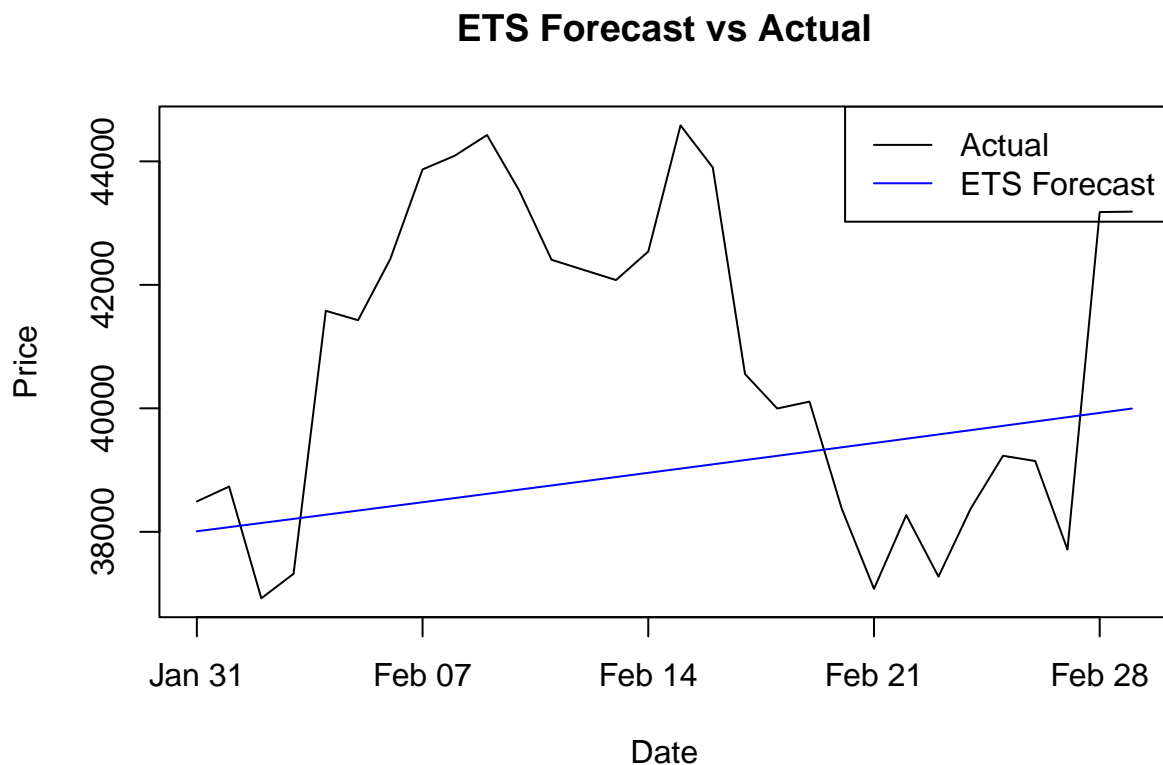
```
## MAE: 2676.604
```

```
cat("RMSE:", round(rmse_ets, 4), "\n")
```

```
## RMSE: 3176.747
```

```
plot
```

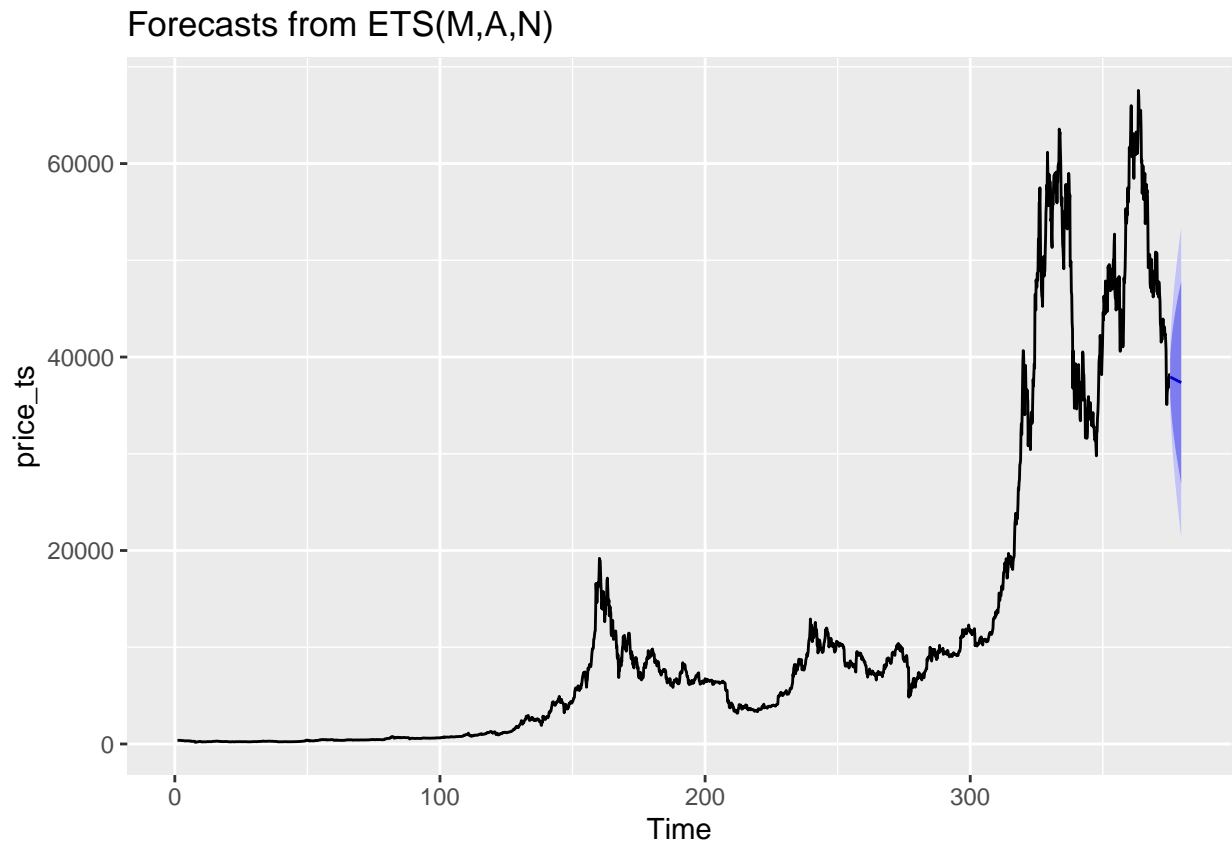
```
plot(index(x)[(length(x) - test_size + 1):length(x)], true_prices,  
      type = "l", col = "black", ylab = "Price", xlab = "Date", main = "ETS Forecast vs Actual")  
lines(index(x)[(length(x) - test_size + 1):length(x)], predicted_prices_ets, col = "blue")  
legend("topright", legend = c("Actual", "ETS Forecast"), col = c("black", "blue"), lty = 1)
```



ETS on just price

```
price_ts <- ts(coredata(x)[1:(length(x) - test_size)], frequency = 7) # weekly freq (adjust as needed)  
  
fit_ets_price <- ets(price_ts)  
fc_price <- forecast(fit_ets_price, h = test_size)  
  
predicted_prices <- fc_price$mean  
true_prices <- coredata(x)[(length(x) - test_size + 1):length(x)]
```

```
ets_raw_pred_price = predicted_prices
autoplot(fc_price)
```



```
mae <- mean(abs(true_prices - predicted_prices))
rmse <- sqrt(mean((true_prices - predicted_prices)^2))

mae
```

```
## [1] 3265.099
```

```
rmse
```

```
## [1] 3982.001
```

```
accuracy(fc_price, true_prices)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1.856009 783.7749 333.4589 -0.006612916 2.601616 0.996717
## Test set     3119.879498 3982.0007 3265.0990 7.308372829 7.699924 9.759462
##           ACF1
## Training set 0.04045423
## Test set     NA
```

GARCH

```
library(rugarch)

best_aic <- Inf
best_model <- NULL
best_spec <-

train_vec <- as.numeric(train)

for (p in 0:2) {
  for (q in 0:2) {
    for (r in 1:2) {
      for (s in 1:2) {
        cat(sprintf("Trying ARMA(%d,%d)-GARCH(%d,%d)\n", p, q, r, s))

        spec <- ugarchspec(
          mean.model = list(armaOrder = c(p, q), include.mean = TRUE),
          variance.model = list(garchOrder = c(r, s)),
          distribution.model = "norm" # Or "std" for Student-t
        )

        tryCatch({
          fit <- ugarchfit(spec, data = train_vec, solver = "hybrid")
          aic <- infocriteria(fit)[1] # AIC

          if (aic < best_aic) {
            best_aic <- aic
            best_model <- fit
            best_spec <- spec
          }
        }, error = function(e) {
          cat("Model failed: ", e$message, "\n")
        })
      }
    }
  }
}
```

```
## Trying ARMA(0,0)-GARCH(1,1)
## Trying ARMA(0,0)-GARCH(1,2)
## Trying ARMA(0,0)-GARCH(2,1)
## Trying ARMA(0,0)-GARCH(2,2)
## Trying ARMA(0,1)-GARCH(1,1)
## Trying ARMA(0,1)-GARCH(1,2)
## Trying ARMA(0,1)-GARCH(2,1)
## Trying ARMA(0,1)-GARCH(2,2)
## Trying ARMA(0,2)-GARCH(1,1)
## Trying ARMA(0,2)-GARCH(1,2)
## Trying ARMA(0,2)-GARCH(2,1)
## Trying ARMA(0,2)-GARCH(2,2)
## Trying ARMA(1,0)-GARCH(1,1)
```

```
## Trying ARMA(1,0)-GARCH(1,2)
## Trying ARMA(1,0)-GARCH(2,1)
## Trying ARMA(1,0)-GARCH(2,2)
## Trying ARMA(1,1)-GARCH(1,1)
## Trying ARMA(1,1)-GARCH(1,2)
## Trying ARMA(1,1)-GARCH(2,1)
## Trying ARMA(1,1)-GARCH(2,2)
## Trying ARMA(1,2)-GARCH(1,1)
## Trying ARMA(1,2)-GARCH(1,2)
## Trying ARMA(1,2)-GARCH(2,1)
## Trying ARMA(1,2)-GARCH(2,2)
## Trying ARMA(2,0)-GARCH(1,1)
## Trying ARMA(2,0)-GARCH(1,2)
## Trying ARMA(2,0)-GARCH(2,1)
## Trying ARMA(2,0)-GARCH(2,2)
## Trying ARMA(2,1)-GARCH(1,1)
## Trying ARMA(2,1)-GARCH(1,2)
## Trying ARMA(2,1)-GARCH(2,1)
## Trying ARMA(2,1)-GARCH(2,2)
## Trying ARMA(2,2)-GARCH(1,1)
## Trying ARMA(2,2)-GARCH(1,2)
## Trying ARMA(2,2)-GARCH(2,1)
## Trying ARMA(2,2)-GARCH(2,2)
```

```
cat("\nBest model AIC:", best_aic, "\n")
```

```
##
## Best model AIC: -3.773393
```

```
show(best_model)
```

```
##
## *-----*
## *          GARCH Model Fit          *
## *-----*
##
## Conditional Variance Dynamics
## -----
## GARCH Model   : sGARCH(1,1)
## Mean Model    : ARFIMA(2,0,2)
## Distribution   : norm
##
## Optimal Parameters
## -----
##      Estimate  Std. Error  t value Pr(>|t|)
## mu      0.001994   0.000661   3.0165 0.002557
## ar1     1.575740   0.021763  72.4039 0.000000
## ar2    -0.849327   0.020577 -41.2758 0.000000
## ma1    -1.590459   0.024182 -65.7713 0.000000
## ma2     0.879828   0.015622  56.3214 0.000000
## omega    0.000062   0.000010   6.3333 0.000000
## alpha1   0.139252   0.015665   8.8894 0.000000
## beta1    0.838630   0.015208  55.1447 0.000000
```

```

##
## Robust Standard Errors:
##      Estimate   Std. Error   t value Pr(>|t|)
## mu      0.001994    0.000689    2.8921 0.003827
## ar1     1.575740    0.024820   63.4874 0.000000
## ar2     -0.849327    0.033927  -25.0338 0.000000
## ma1     -1.590459    0.015873 -100.1975 0.000000
## ma2      0.879828    0.021770   40.4146 0.000000
## omega    0.000062    0.000025    2.5406 0.011066
## alpha1   0.139252    0.037162    3.7471 0.000179
## beta1    0.838630    0.027188   30.8451 0.000000
##
## LogLikelihood : 4951.145
##
## Information Criteria
## -----
##
## Akaike          -3.7734
## Bayes           -3.7555
## Shibata         -3.7734
## Hannan-Quinn   -3.7669
##
## Weighted Ljung-Box Test on Standardized Residuals
## -----
##
##              statistic   p-value
## Lag[1]              3.822 0.0505840
## Lag[2*(p+q)+(p+q)-1][11]  8.266 0.0003371
## Lag[4*(p+q)+(p+q)-1][19] 15.062 0.0274064
## d.o.f=4
## H0 : No serial correlation
##
## Weighted Ljung-Box Test on Standardized Squared Residuals
## -----
##
##              statistic   p-value
## Lag[1]              0.234 0.6286
## Lag[2*(p+q)+(p+q)-1][5]  2.095 0.5961
## Lag[4*(p+q)+(p+q)-1][9]  3.111 0.7404
## d.o.f=2
##
## Weighted ARCH LM Tests
## -----
##
##      Statistic Shape Scale P-Value
## ARCH Lag[3]      1.725 0.500 2.000 0.1890
## ARCH Lag[5]      2.865 1.440 1.667 0.3100
## ARCH Lag[7]      3.249 2.315 1.543 0.4678
##
## Nyblom stability test
## -----
## Joint Statistic: 1.9205
## Individual Statistics:
## mu      0.12078
## ar1     0.03063
## ar2     0.06082
## ma1     0.06063

```



```
## ma2      0.05446
## omega    0.70892
## alpha1   0.08441
## beta1    0.41949
##
## Asymptotic Critical Values (10% 5% 1%)
## Joint Statistic:      1.89 2.11 2.59
## Individual Statistic: 0.35 0.47 0.75
##
## Sign Bias Test
## -----
##              t-value   prob sig
## Sign Bias      1.324 0.1856
## Negative Sign Bias 1.063 0.2880
## Positive Sign Bias 1.417 0.1565
## Joint Effect    3.191 0.3631
##
##
## Adjusted Pearson Goodness-of-Fit Test:
## -----
##   group statistic p-value(g-1)
## 1    20      426.9   1.120e-78
## 2    30      445.5   4.131e-76
## 3    40      465.6   1.882e-74
## 4    50      464.5   4.870e-69
##
##
## Elapsed time : 0.3591509
```

```
best_garch = best_model
```

The model is Conditional Variance Dynamics, GARCH Model : sGARCH(1,1), Mean Model :ARFIMA(2,0,2), Distribution : norm

```
# garch_spec <- ugarchspec(
#   mean.model      = list(armaOrder = c(1, 1), include.mean = TRUE),
#   variance.model = list(garchOrder = c(1, 1)),
#   distribution.model = "norm" # you can change to "std" for Student-t
# )
#
#
# # train is your zoo object of log-returns, from earlier
# train_vec <- as.numeric(train) # convert to numeric vector
#
# garch_fit <- ugarchfit(spec = garch_spec, data = train_vec)
# show(garch_fit)
```

```
# use best one
garch_fc <- ugarchforecast(best_garch, n.ahead = test_size)

# Extract forecasted log-returns
fc_log_returns_garch <- fitted(garch_fc)
```

```

last_price <- coredata(x)[length(x) - test_size]

# Reconstruct forecasted prices from log-returns
predicted_prices_garch <- numeric(test_size)
predicted_prices_garch[1] <- last_price * exp(fc_log_returns_garch[1])

for (i in 2:test_size) {
  predicted_prices_garch[i] <- predicted_prices_garch[i - 1] * exp(fc_log_returns_garch[i])
}

true_prices <- coredata(x)[(length(x) - test_size + 1):length(x)]

mae_garch <- mean(abs(true_prices - predicted_prices_garch))
rmse_garch <- sqrt(mean((true_prices - predicted_prices_garch)^2))

cat("GARCH Forecast\n")

## GARCH Forecast

cat("MAE : ", round(mae_garch, 4), "\n")


## MAE : 2688.547

cat("RMSE: ", round(rmse_garch, 4), "\n")

## RMSE: 3170.944

plot(fc_log_returns_garch)

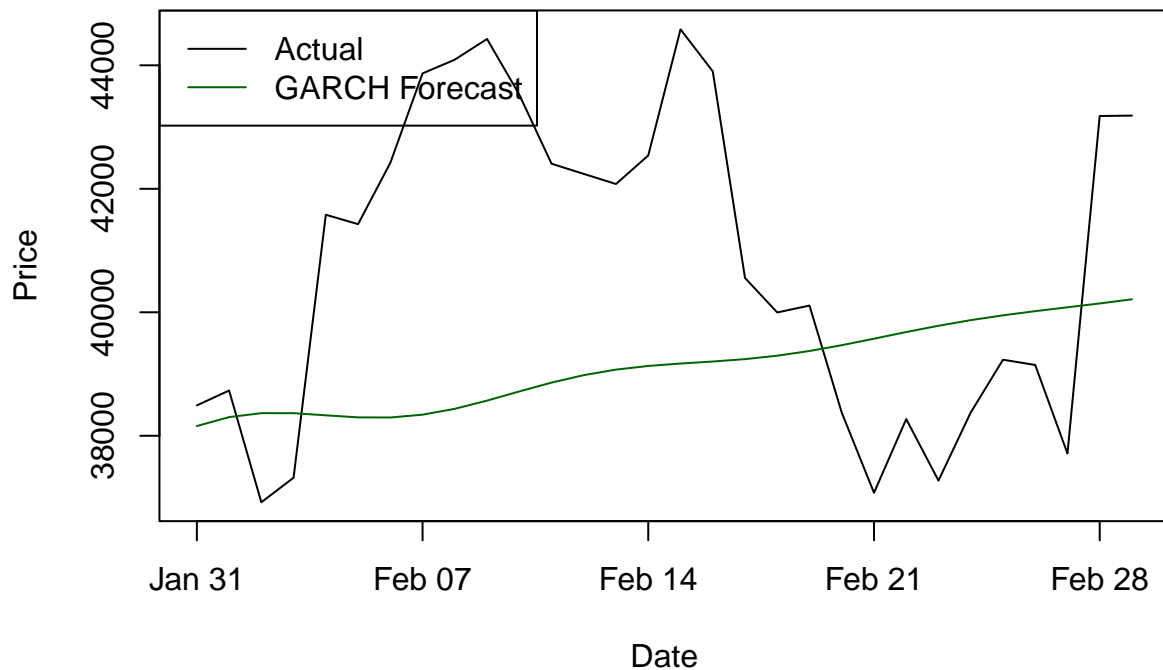
```



plot

```
plot(index(x)[(length(x) - test_size + 1):length(x)], true_prices,
     type = "l", col = "black", ylab = "Price", xlab = "Date", main = "GARCH Forecast vs Actual")
lines(index(x)[(length(x) - test_size + 1):length(x)], predicted_prices_garch, col = "darkgreen")
legend("topleft", legend = c("Actual", "GARCH Forecast"), col = c("black", "darkgreen"), lty = 1)
```

GARCH Forecast vs Actual



Summary

Regression (Forecasting)

```
evaluate_model <- function(model_type = c("ARIMA", "ETS", "GARCH"),
                           train, test, full_price_series, test_size = 30, best_garch) {

  model_type <- match.arg(model_type)

  # Reconstruct price from forecasted log-returns
  reconstruct_prices <- function(log_returns, last_price) {
    pred <- numeric(length(log_returns))
    pred[1] <- last_price * exp(log_returns[1])
    for (i in 2:length(log_returns)) {
      pred[i] <- pred[i - 1] * exp(log_returns[i])
    }
    pred
  }

  last_price <- coredata(full_price_series)[length(full_price_series) - test_size]
  true_prices <- coredata(full_price_series)[(length(full_price_series) - test_size + 1):length(full_price_series)]

  if (model_type == "ARIMA") {
```

```

fit <- auto.arima(train)
train_pred <- fitted(fit)
test_fc <- forecast(fit, h = test_size)$mean

} else if (model_type == "ETS") {
  fit <- ets(train)
  train_pred <- fitted(fit)
  test_fc <- forecast(fit, h = test_size)$mean

} else if (model_type == "GARCH") {

  # garch_spec <- ugarchspec(
  #   mean.model = list(armaOrder = c(p, q), include.mean = TRUE),
  #   variance.model = list(garchOrder = c(r, s)),
  #   distribution.model = dist
  # )
  # fit <- ugarchfit(best_garch, as.numeric(train))
  # train_pred <- fitted(fit)
  train_pred <- fitted(best_garch)

  garch_fc <- ugarchforecast(best_garch, n.ahead = test_size)
  test_fc <- fitted(garch_fc)
}

# Price predictions
predicted_train_prices <- reconstruct_prices(train_pred, coredata(full_price_series)[length(full_price_series) - test_size + 1:length(full_price_series)])
predicted_test_prices <- reconstruct_prices(test_fc, last_price)

train_class_pred <- ifelse(train_pred >= 0, 1, 0)
test_class_pred <- ifelse(test_fc >= 0, 1, 0)

train_class_actual <- ifelse(train >= 0, 1, 0)
test_class_actual <- ifelse(test >= 0, 1, 0)

train_acc = Accuracy(train_class_pred, train_class_actual)
test_acc = Accuracy(test_class_actual, test_class_pred)

train_f1 = F1_Score(train_class_actual, train_class_pred)
test_f1 = F1_Score(test_class_actual, test_class_pred)

train_recall = Recall(train_class_actual, train_class_pred)
test_recall = Recall(test_class_actual, test_class_pred)

train_precision = Precision(train_class_actual, train_class_pred)
test_precision = Precision(test_class_actual, test_class_pred)

actual_train_prices <- coredata(full_price_series)[(length(full_price_series) - test_size - length(train_pred) + 1):length(full_price_series)]

train_mae <- mean(abs(actual_train_prices - predicted_train_prices))

```

```

train_rmse <- sqrt(mean((actual_train_prices - predicted_train_prices)^2))
test_mae <- mean(abs(true_prices - predicted_test_prices))
test_rmse <- sqrt(mean((true_prices - predicted_test_prices)^2))

data.frame(
  Model = model_type,
  Train_MAE = round(train_mae, 4),
  Test_MAE = round(test_mae, 4),
  Train_RMSE = round(train_rmse, 4),
  Test_RMSE = round(test_rmse, 4),
  Train_Accuracy = round(train_acc, 4),
  Test_Accuracy = round(test_acc, 4),
  Train_F1 = round(train_f1, 4),
  Test_F1 = round(test_f1, 4),
  Train_Recall = round(train_recall, 4),
  Test_Recall = round(test_recall, 4),
  Train_Precision = round(train_precision, 4),
  Test_Precision = round(test_precision, 4)
)
}

```

```

evaluate_model(model = "GARCH", train = train, test = test, full_price_series = x, best_garch = best_ga

```

```

## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.zoo", "Ops.xts")
## for "=="

```

```

##   Model Train_MAE Test_MAE Train_RMSE Test_RMSE Train_Accuracy Test_Accuracy
## 1 GARCH 4307.186 2688.547 6806.504 3170.944 0.5385 0.4667
##   Train_F1 Test_F1 Train_Recall Test_Recall Train_Precision Test_Precision
## 1 0.3103 0.1111 0.2254 0.0714 0.4982 0.25

```

```

results <- rbind(
  evaluate_model("ARIMA", train, test, x, test_size),
  evaluate_model("ETS", train, test, x, test_size),
  evaluate_model("GARCH", train, test, x, test_size, best_garch = best_garch)
)

```

```

## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.zoo", "Ops.ts")
## for "=="

```

```

## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.ts", "Ops.zoo")
## for "=="

```

```

## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.zoo", "Ops.ts")
## for "=="

```

```

## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.ts", "Ops.zoo")
## for "=="

```

```
## Warning in mean(y_true == y_pred): Incompatible methods ("Ops.zoo", "Ops.xts")
## for "=="
```

```
results
```

```
##   Model Train_MAE Test_MAE Train_RMSE Test_RMSE Train_Accuracy Test_Accuracy
## 1 ARIMA 4880.190 2671.279 9363.207 3167.454 0.5447 0.5333
## 2 ETS 4725.475 2676.604 9185.440 3176.747 0.5393 0.5333
## 3 GARCH 4307.186 2688.547 6806.504 3170.944 0.5385 0.4667
##   Train_F1 Test_F1 Train_Recall Test_Recall Train_Precision Test_Precision
## 1 0.2308 NA 0.1483 NA 0.5203 NA
## 2 NA NA NA NA NA NA
## 3 0.3103 0.1111 0.2254 0.0714 0.4982 0.25
```

```
# write.csv(results, file = "figs/forecast_results_classical_methods.csv", row.names = FALSE)
```

Classification

Done in above.